

# Amazon SQS Integration for DriversKlub Backend

## Executive Summary

This document outlines the proposal to integrate **Amazon Simple Queue Service (SQS)** as the message queue system for the DriversKlub backend microservices architecture. SQS will enable asynchronous, event-driven communication between services, improving scalability, reliability, and system decoupling.

## Why Amazon SQS?

### Benefits Over Kafka

Feature	Amazon SQS	Apache Kafka
Deployment Complexity	✓ Zero - Fully managed	✗ High - Requires infrastructure
Maintenance	✓ None - AWS managed	✗ Ongoing - Self-managed
Cost (Low Volume)	✓ Free tier (1M requests/month)	✗ EC2/EKS costs even at low volume
Scalability	✓ Automatic	⚠ Manual configuration
Learning Curve	✓ Simple API	✗ Steep learning curve
Reliability	✓ 99.9% SLA, built-in redundancy	⚠ Depends on setup
Setup Time	✓ Minutes	✗ Days/Weeks

### Key Advantages

#### 1. Zero Infrastructure Management

- No servers to provision or maintain
- No capacity planning required
- Automatic scaling based on load

#### 2. Cost-Effective

- **Free Tier:** 1 million requests/month
- **Pay-as-you-go:** \$0.40 per million requests after free tier
- **Estimated cost:** ~\$1-2/month for current scale

#### 3. High Reliability

- 99.9% availability SLA
- Automatic message replication across multiple availability zones
- Built-in Dead Letter Queue (DLQ) for failed messages

#### 4. Developer Friendly

- Simple REST API
- Official AWS SDK for Node.js

- Extensive documentation and community support

## 5. Production Ready

- Used by thousands of companies globally
  - Battle-tested at massive scale
  - Integrated with AWS ecosystem (CloudWatch, IAM, etc.)
- 

# Use Cases in DriversKlub

---

## 1. Trip Lifecycle Events

**Problem:** When a trip is created, multiple services need to be notified (notifications, analytics, provider sync).

**Solution:** Publish `trip.created` event to SQS queue. Multiple consumers process it independently.

```
Trip Service → [SQS Queue] → Notification Service (sends SMS)
                           → Analytics Service (logs event)
                           → Provider Service (syncs with MojoBoxx/MMT)
```

### Benefits:

- Trip Service doesn't wait for notifications to be sent
- If notification fails, trip creation still succeeds
- Can add new consumers without modifying Trip Service

## 2. Driver Status Changes

**Problem:** When a driver goes online/offline, multiple systems need updates (assignment service, analytics, Rapido sync).

**Solution:** Publish `driver.status_changed` event.

```
Driver Service → [SQS Queue] → Assignment Service (updates availability)
                           → Rapido Service (syncs captain status)
                           → Analytics Service (tracks uptime)
```

## 3. Payment Processing

**Problem:** Payment processing can be slow and should not block trip completion.

**Solution:** Use FIFO queue for guaranteed order and exactly-once processing.

```
Trip Service → [SQS FIFO Queue] → Payment Service (processes payment)
                           → Accounting Service (records transaction)
```

## 4. Provider Webhooks

**Problem:** External providers (MojoBoxx, MMT, Rapido) send webhooks that need reliable processing.

**Solution:** Queue webhook payloads for asynchronous processing.

```
Webhook Endpoint → [SQS Queue] → Worker Service (processes webhook)
→ Retry on failure (up to 3 times)
→ DLQ if still failing
```

## 5. Notification Delivery

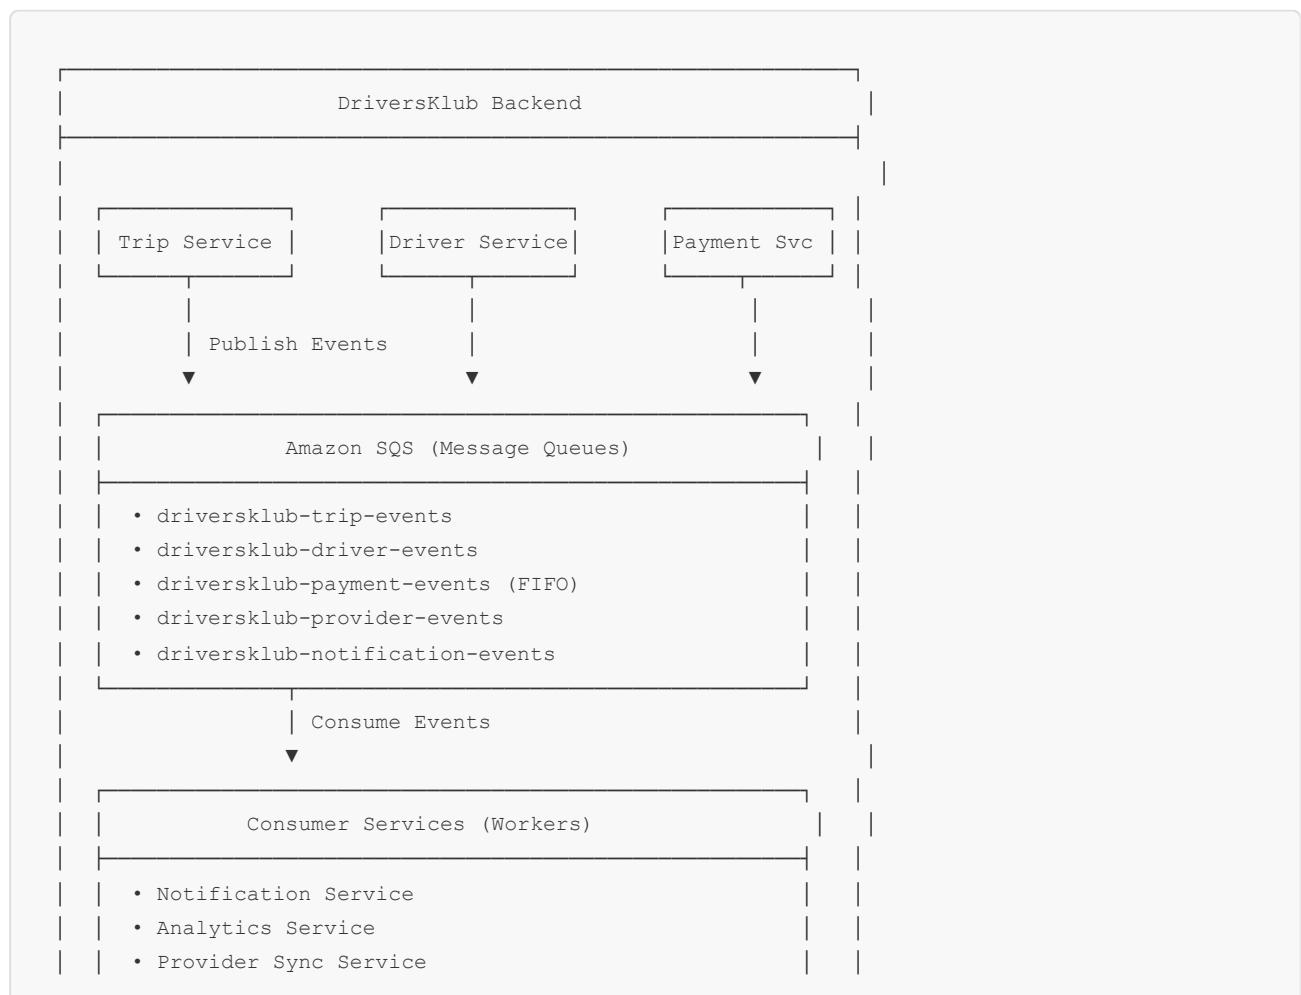
**Problem:** Sending SMS/Push notifications can fail or be slow.

**Solution:** Queue all notifications for async delivery with retry logic.

```
Any Service → [SQS Queue] → Notification Service → SMS Gateway
→ Push Notification Service
→ Email Service
```

# Architecture Overview

## Queue Structure





## Event Flow Example: Trip Creation

1. User books a trip via API  
↓
2. Trip Service creates trip in database  
↓
3. Trip Service publishes "trip.created" event to SQS  
↓
4. Trip Service returns success to user (fast response)  
↓
5. SQS delivers message to multiple consumers:
  - Notification Service → Sends booking confirmation SMS
  - Provider Service → Books with MojoBoxx/MMT
  - Analytics Service → Logs trip metrics↓
6. Each consumer processes independently
  - If one fails, others still succeed
  - Failed messages go to Dead Letter Queue for investigation

## Technical Implementation

---

### Queue Types

#### Standard Queue (Default)

- **Use for:** Most events (trips, drivers, notifications)
- **Delivery:** At-least-once (may receive duplicates)
- **Ordering:** Best-effort (not guaranteed)
- **Throughput:** Nearly unlimited
- **Cost:** \$0.40 per million requests

#### FIFO Queue

- **Use for:** Payments, critical operations
- **Delivery:** Exactly-once (no duplicates)
- **Ordering:** Guaranteed order
- **Throughput:** 300 transactions/second
- **Cost:** \$0.50 per million requests

### Event Schema Example

```
// Trip Created Event
{
```

```

    "eventType": "trip.created",
    "eventId": "evt_123456",
    "timestamp": "2026-01-13T13:00:00Z",
    "data": {
        "tripId": "trip_abc123",
        "userId": "user_xyz789",
        "pickupLocation": {
            "lat": 28.6139,
            "lng": 77.2090,
            "address": "Connaught Place, New Delhi"
        },
        "dropLocation": {
            "lat": 28.5355,
            "lng": 77.3910,
            "address": "Noida Sector 62"
        },
        "scheduledTime": "2026-01-13T15:00:00Z",
        "vehicleType": "SEDAN",
        "fare": 1500
    }
}

```

## Dead Letter Queue (DLQ)

**Purpose:** Capture messages that fail processing after multiple retries.

### Configuration:

- Max retries: 3 attempts
- Retry delay: Exponential backoff (1s, 2s, 4s)
- DLQ retention: 14 days

**Monitoring:** CloudWatch alarms when DLQ receives messages.

## Cost Analysis

### Assumptions

- **Trips per day:** 10,000
- **Events per trip:** 5 (created, assigned, started, completed, payment)
- **Total events/day:** 50,000
- **Monthly events:** 1.5 million

### Cost Breakdown

Component	Volume	Unit Cost	Monthly Cost
Standard Queues	1.5M requests	FREE (under 1M) + \$0.40/M	\$0.20
FIFO Queue (Payments)	300K requests	\$0.50/M	\$0.15

Data Transfer	Negligible	\$0.09/GB	<b>\$0.05</b>
<b>TOTAL</b>			<b>\$0.40/month</b>

**Note:** First 1M requests are FREE, so actual cost is even lower initially.

## Comparison with Kafka

Component	SQS	Self-Hosted Kafka
Infrastructure	\$0	\$50-100/month (EC2)
Maintenance	\$0	10-20 hours/month
Monitoring	Included	\$20/month (tools)
<b>Total</b>	<b>\$0.40/month</b>	<b>\$70-120/month</b>

**Savings:** ~\$70-120/month + significant developer time

---

## Implementation Roadmap

---

### Phase 1: Setup & Infrastructure (Week 1)

**Duration:** 2-3 days

- Create AWS account (if needed)
- Set up IAM user with SQS permissions
- Create SQS queues via AWS Console or script
- Install AWS SDK (`@aws-sdk/client-sqs`)
- Create shared queue package (`@driversklub/queue`)
- Configure environment variables

**Deliverables:**

- All queues created in AWS
- Queue URLs documented
- Basic publish/consume example working

### Phase 2: Core Integration (Week 2)

**Duration:** 5-7 days

- Implement queue manager (publish/consume utilities)
- Define event schemas with TypeScript types
- Add event publishing to Trip Service
- Add event publishing to Driver Service
- Create consumer for Notification Service
- Implement retry logic and DLQ handling

**Deliverables:**

- Trip and Driver events publishing successfully
- Notifications being sent via queue
- DLQ monitoring in place

**Phase 3: Extended Features (Week 3)****Duration:** 5-7 days

- Add payment event processing (FIFO queue)
- Implement provider webhook queuing
- Add analytics event consumers
- Create monitoring dashboard
- Write integration tests
- Performance testing

**Deliverables:**

- All event types implemented
- Comprehensive test coverage
- Monitoring and alerting configured

**Phase 4: Production Deployment (Week 4)****Duration:** 3-5 days

- Deploy to staging environment
- Load testing with production-like data
- Monitor CloudWatch metrics
- Gradual rollout to production
- Documentation and runbooks

**Deliverables:**

- Production deployment complete
- Team training completed
- Operational runbooks ready

**Total Timeline:** 3-4 weeks for full implementation

---

**Minimal Viable Product (MVP) - Quick Start**

---

For immediate testing and validation, we can implement a **minimal version in 1-2 days**:**MVP Scope**

1. **Single Queue:** `driversklub-events` (Standard Queue)
2. **One Publisher:** Trip Service publishes `trip.created` events
3. **One Consumer:** Notification Service sends SMS on trip creation
4. **Basic Error Handling:** DLQ for failed messages

## MVP Benefits

- Prove the concept works
- Test AWS integration
- Validate event schema design
- Measure performance
- Get team familiar with SQS

## MVP to Full Implementation

Once MVP is validated, expand incrementally:

- Add more event types
  - Add more consumers
  - Implement FIFO queues for payments
  - Add monitoring and alerting
- 

## Risk Assessment

---

### Technical Risks

Risk	Impact	Probability	Mitigation
AWS account issues	High	Low	Set up IAM properly, test credentials
Message loss	High	Very Low	SQS has 99.9% SLA, use DLQ
Duplicate processing	Medium	Low	Implement idempotency keys
Cost overrun	Low	Very Low	Set CloudWatch billing alarms
Learning curve	Low	Medium	Comprehensive documentation, training

### Operational Risks

Risk	Impact	Probability	Mitigation
DLQ fills up	Medium	Low	CloudWatch alarms, daily monitoring
Consumer lag	Medium	Low	Auto-scaling consumers, monitoring
AWS region outage	High	Very Low	Multi-region setup (future)

**Overall Risk Level: LOW** - SQS is a mature, battle-tested service

---

## Monitoring & Observability

---

## CloudWatch Metrics (Built-in)

- **ApproximateNumberOfMessagesVisible**: Messages in queue
- **ApproximateNumberOfMessagesNotVisible**: Messages being processed
- **NumberOfMessagesSent**: Publishing rate
- **NumberOfMessagesReceived**: Consumption rate
- **ApproximateAgeOfOldestMessage**: Queue lag indicator

## Custom Metrics

- Event processing time
- Success/failure rates per event type
- DLQ message count by error type

## Alerts

- DLQ receives messages → Slack/Email alert
- Queue depth > 1000 → Scale up consumers
- Processing time > 30s → Investigate performance

# Security Considerations

## IAM Permissions

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "sns:Publish",
                "sns:DeleteTopic",
                "sns:ListTopics"
            ],
            "Resource": "arn:aws:sns:us-east-1:*:driversclub-*"
        }
    ]
}
```

## Best Practices

- Use IAM roles (not hardcoded keys) in production
- Encrypt messages at rest (AWS KMS)
- Encrypt messages in transit (HTTPS)
- Separate IAM users for dev/staging/prod
- Rotate credentials regularly
- Use VPC endpoints for private communication

## Success Metrics

---

### Technical Metrics

- **Message Processing Time:** < 5 seconds (p95)
- **Queue Depth:** < 100 messages during normal operation
- **DLQ Messages:** < 0.1% of total messages
- **Availability:** > 99.9%

### Business Metrics

- **Notification Delivery:** > 99% success rate
- **Trip Processing:** 100% reliability
- **System Decoupling:** Services can deploy independently
- **Developer Velocity:** Faster feature development

---

## Alternatives Considered

---

### 1. Apache Kafka

**Pros:** High throughput, event streaming, complex event processing **Cons:** Complex setup, high operational overhead, overkill for current scale **Verdict:** ❌ Too complex for current needs

### 2. RabbitMQ

**Pros:** Feature-rich, good for complex routing **Cons:** Self-hosted, requires maintenance, learning curve **Verdict:** ❌ Still requires infrastructure management

### 3. Redis Pub/Sub

**Pros:** Fast, simple, already using Redis **Cons:** No persistence, no guaranteed delivery, not designed for queuing **Verdict:** ❌ Not reliable enough for critical events

### 4. Database-based Queue

**Pros:** No new infrastructure, simple **Cons:** Poor performance at scale, polling overhead, not designed for queuing **Verdict:** ❌ Not scalable

### 5. Amazon SQS

**Pros:** Fully managed, reliable, cost-effective, simple, scalable **Cons:** AWS vendor lock-in (minor concern) **Verdict:**  **Best fit for current requirements**

---

## Conclusion

---

Amazon SQS provides the optimal balance of:

-  **Simplicity:** Minimal setup and maintenance
-  **Reliability:** 99.9% SLA with built-in redundancy

- **Cost:** Free tier + pay-as-you-go pricing
- **Scalability:** Automatic scaling with load
- **Developer Experience:** Simple API, great documentation

## Recommendation

**Proceed with Amazon SQS integration** using the phased approach:

1. **Week 1:** MVP implementation (1 queue, 1 publisher, 1 consumer)
2. **Week 2-3:** Expand to all event types and consumers
3. **Week 4:** Production deployment with monitoring

## Next Steps

1. **Approve this proposal**
  2. **Set up AWS account** (if not already done)
  3. **Create IAM credentials** for development
  4. **Begin MVP implementation** (estimated 1-2 days)
- 

## Appendix

---

### A. Estimated Development Effort

Task	Effort (Hours)
AWS setup and queue creation	4
Queue package development	8
Event schema definition	4
Trip Service integration	6
Driver Service integration	6
Notification consumer	6
Payment consumer (FIFO)	8
Testing and debugging	12
Documentation	4
<b>Total</b>	<b>58 hours (~1.5 weeks)</b>

### B. Required AWS Permissions

```

  sqs:CreateQueue
  sqs:DeleteQueue
  sqs:GetQueueAttributes
  sqs:SetQueueAttributes
  sqs:SendMessage

```

```
  sqs:ReceiveMessage  
  sqs:DeleteMessage  
  sqs:PurgeQueue  
  cloudwatch:PutMetricData
```

## C. Useful Resources

- [AWS SQS Documentation](#)
- [AWS SDK for JavaScript](#)
- [SQS Best Practices](#)
- [SQS Pricing Calculator](#)

---

**Document Version:** 1.0

**Date:** January 13, 2026

**Author:** DriversKlub Engineering Team

**Status:** Pending Approval